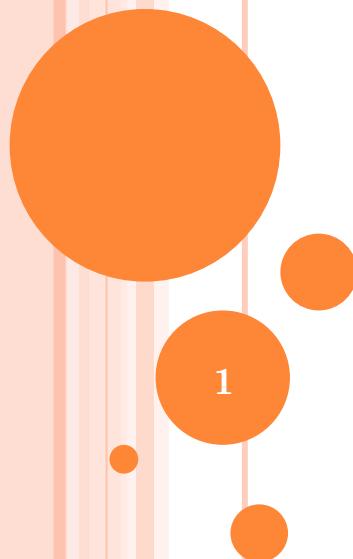




CSE 4392 SPECIAL TOPICS
NATURAL LANGUAGE PROCESSING

Log-linear Models



2025 Spring

LAST TIME

- Supervised classification:
 - Document to classify, d
 - Set of classes, $C = \{c_1, c_2, \dots, c_k\}$
- Naive Bayes:

$$\hat{c} = \arg \max_c P(c)P(d|c)$$

LOGISTIC REGRESSION

- Powerful supervised model
- Baseline approach to most NLP tasks
- Connections with neural networks
- Binary (two classes) or multinomial (>2 classes)

DISCRIMINATIVE MODEL

- Logistic Regression is a *discriminative* model
- Naive Bayes is a *generative* model



DISCRIMINATIVE MODEL

- Logistic Regression:

$$\hat{c} = \arg \max_{c \in C} P(c|d)$$

- Naive Bayes:

$$\hat{c} = \arg \max_{c \in C} P(c) P(d|c)$$



QUIZ



- Given that we want to classify an image into either a dog or a cat (no other choices), name the features you would use (can be numerical or categorical).

USING LOGISTIC REGRESSION

- Inputs:

1. Classification instance in a **feature representation** $[x_1, x_2, \dots, x_d]$
2. **Classification function** to compute \hat{y} using $P(\hat{y} | x)$
3. **Loss function** (for learning)
4. Optimization **algorithm**

- Train phase:

- Learn the **parameters** of the model to minimize **loss function**

- Test phase:

- Apply **parameters** to predict class given a new input x

FEATURE REPRESENTATION

- Input observation: $x^{(i)}$
- Feature vector: $[x_1, x_2, \dots, x_d]$
- Feature j of i^{th} input: $x_j^{(i)}$

SAMPLE FEATURE VECTOR

It's hokey. There are virtually no surprises, and the writing is second-rate.
So why was it so enjoyable? For one thing, the cast is great. Another nice touch is the music I was overcome with the urge to get off the couch and start dancing. It sucked me in, and it'll do the same to you.

$$x_2 = 2$$

$$x_3 = 1$$

$$x_1 = 3$$

$$x_5 = 0$$

$$x_6 = 4.15$$


$$x_4 = 3$$

Var	Definition	Value
x_1	count(positive lexicon) \in doc)	3
x_2	count(negative lexicon) \in doc)	2
x_3	$\begin{cases} 1 & \text{if "no" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
x_4	count(1st and 2nd pronouns \in doc)	3
x_5	$\begin{cases} 1 & \text{if "!" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
x_6	log(word count of doc)	$\ln(64) = 4.15$

CLASSIFICATION FUNCTION

- Given: Input feature vector $[x_1, x_2, \dots, x_d]$
- Output: $P(y = 1 | x)$ and $P(y = 0 | x)$ *(binary classification)*
- Require a function, $F : \mathbb{R}^d \rightarrow [0,1]$ *(probability)*
- Sigmoid (or logistic) Function:

$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + \exp(-z)}$$



QUIZ

- Why do we use Sigmoid/Logistic function as our classification function? (Select all that apply)
 - a) Produces a value between 0 and 1
 - b) A partial function with domain $[0, +\infty)$
 - c) Produces a value between -1 to 1
 - d) A total function with domain $(-\infty, \infty)$
 - e) Integrates to 1 from $-\infty$ to ∞
 - f) Differentiable

WEIGHTS AND BIASES

- *Which features are important and how much?*
- Learn a vector of **weights** and a **bias**
- **Weights:** Vector of real numbers,

$$w = [w_1, w_2, \dots, w_d]$$

- **Bias:** Scalar intercept, b
- Given an instance \mathbf{x} :

$$z = \sum_{i=1}^d w_i x_i + b$$

$$\text{or } z = \mathbf{w} \cdot \mathbf{x} + b$$

WHAT IS THE BIAS?

$$z = \mathbf{w} \cdot \mathbf{x} + b$$

- Bias, or intercept, gives the default behavior of the classifier when no useful information about \mathbf{x} is known.
- Try setting x_i to be all 0:

$$z = b$$

- Gives the prior probability distribution of the classes without looking at the input features:

prediction_bias = avg_predictions - avg of labels in data set

PUTTING IT TOGETHER

- Given \mathbf{x} , compute $z = \mathbf{w} \cdot \mathbf{x} + b$
- Compute probabilities:


$$\begin{aligned} P(y = 1 | \mathbf{x}) &= \frac{1}{1 + e^{-z}} \\ P(y = 1) &= \sigma(\mathbf{w} \cdot \mathbf{x} + b) \\ &= \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}} \\ P(y = 0) &= 1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b) \\ &= 1 - \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}} \\ &= \frac{e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}} \end{aligned}$$

- Decision boundary:

$$\hat{y} = \begin{cases} 1 & \text{if } P(y = 1 | \mathbf{x}) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

PUTTING IT TOGETHER


- Given \mathbf{x} , compute $z = \mathbf{w} \cdot \mathbf{x} + b$
- Compute probabilities:



- Decision boundary:

$$\hat{y} = \begin{cases} 1 & \text{if } P(y = 1 | x) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

EXAMPLE: SENTIMENT CLASSIFICATION



Var	Definition	Value
x_1	count(positive lexicon) \in doc	3
x_2	count(negative lexicon) \in doc)	2
x_3	$\begin{cases} 1 & \text{if "no" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
x_4	count(1st and 2nd pronouns \in doc)	3
x_5	$\begin{cases} 1 & \text{if "!" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
x_6	$\ln(\text{word count of doc})$	$\ln(64) = 4.15$

EXAMPLE: SENTIMENT CLASSIFICATION

Var	Definition	Value
x_1	count(positive lexicon) \in doc)	3
x_2	count(negative lexicon) \in doc)	2
x_3	$\begin{cases} 1 & \text{if "no" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
x_4	count(1st and 2nd pronouns \in doc)	3
x_5	$\begin{cases} 1 & \text{if "!" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
x_6	log(word count of doc)	$\ln(64) = 4.15$

- Assume weights $\mathbf{w} = [2.5, -5.0, -1.2, 0.5, 2.0, 0.7]$ and bias $b = 0.1$

$$\begin{aligned}
 p(+|x) = P(Y = 1|x) &= \sigma(w \cdot x + b) \\
 &= \sigma([2.5, -5.0, -1.2, 0.5, 2.0, 0.7] \cdot [3, 2, 1, 3, 0, 4.15] + 0.1) \\
 &= \sigma(.805) \\
 &= 0.69
 \end{aligned}$$

$$\begin{aligned}
 p(-|x) = P(Y = 0|x) &= 1 - \sigma(w \cdot x + b) \\
 &= 0.31
 \end{aligned}$$

FEATURE DESIGN

- Most important rule: Data is *key*!
- Linguistic intuition (e.g. part of speech tags, parse trees)
- Complex combinations

$$x_1 = \begin{cases} 1 & \text{if } \text{"Case}(w_i) = \text{Lower"} \\ 0 & \text{otherwise} \end{cases}$$

$$x_2 = \begin{cases} 1 & \text{if } w_i \in \text{AcronymDict} \\ 0 & \text{otherwise} \end{cases}$$

$$x_3 = \begin{cases} 1 & \text{if } w_i = \text{St.} \& \text{Case}(w_{i-1}) = \text{Cap} \\ 0 & \text{otherwise} \end{cases}$$

- Feature templates
 - Sparse representations, hash only seen features into index
 - Ex. Trigram("logistic regression model") = Feature #78
- Advanced: Representation learning (we will see this later!)

PROS AND CONS OF LOGISTIC REGRESSION

- More freedom in designing features
 - No strong independence assumptions like Naive Bayes
 - More robust to correlated features (“San Francisco” vs “Boston”) —LR is likely to work better than NB
 - Can even have the same feature twice! (*why?*)
- **However:** Naïve Bayes (NB) often better on very small datasets

LEARNING

- We have our **classification function** - how to assign weights and bias?
- **Goal:** predicted label \hat{y} as close as possible to actual label y
 - **Distance metric/Loss function** between \hat{y} and y :
$$L(\hat{y}, y)$$
 - **Optimization algorithm** for updating weights

LOSS FUNCTION

- Assume $\hat{y} = \sigma(\mathbf{w} \cdot \mathbf{x} + b)$
- $L(\hat{y}, y)$ = Measure of difference between \hat{y} and y . But what form?
- Maximum likelihood estimation (conditional):
 - Choose w and b such that $\log P(y \mid \mathbf{x})$ is maximized for true labels y paired with input x
 - Similar to language models!
 - $\max \log P(w_t \mid w_{t-n}, \dots, w_{t-1})$ given a corpus

CROSS ENTROPY LOSS

- Assume a single data point (x, y) and two classes
- Binary classifier probability (Bernoulli distribution):

$$P(y | x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

- Remember $0 \leq \hat{y} \leq 1$.
- Log probability:
$$\begin{aligned}\log P(y|x) &= \log[\hat{y}^y (1 - \hat{y})^{1-y}] \\ &= y \log \hat{y} + (1 - y) \log(1 - \hat{y})\end{aligned}$$
- CE Loss (we want to minimize):
$$\begin{aligned}-\log P(y|x) &= -y \log \hat{y} - (1 - y) \log(1 - \hat{y}) \\ &= \begin{cases} -\log \hat{y} & \text{if } y = 1 \\ -\log(1 - \hat{y}) & \text{if } y = 0 \end{cases}\end{aligned}$$

CROSS ENTROPY LOSS

- Assume n data points $(x^{(i)}, y^{(i)})$

- Classifier probability:

$$\prod_{i=1}^n P(y|x) = \prod_{i=1}^n \hat{y}^y (1 - \hat{y})^{1-y}$$

(I omitted the (i) here for brevity)

- CE Loss:

$$\begin{aligned} L_{CE} &= -\log \prod_{i=1}^n P(y^{(i)}|x^{(i)}) \\ &= -\sum_{i=1}^n [y^{(i)} \log \hat{y}^{(i)} - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] \end{aligned}$$

EXAMPLE: COMPUTING CE LOSS

Var	Definition	Value
x_1	$\text{count}(\text{positive lexicon}) \in \text{doc}$	3
x_2	$\text{count}(\text{negative lexicon}) \in \text{doc}$	2
x_3	$\begin{cases} 1 & \text{if "no" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
x_4	$\text{count}(1\text{st and 2nd pronouns } \in \text{doc})$	3
x_5	$\begin{cases} 1 & \text{if "!" } \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
x_6	$\log(\text{word count of doc})$	$\ln(64) = 4.15$

- Assume weights $w = [2.5, -5.0, -1.2, 0.5, 2.0, 0.7]$ and bias $b = 0.1$
- If $y = 1$ (positive sentiment), $\text{LCE} = -\log(0.69) = 0.37$
- If $y = 0$ (negative sentiment), $\text{LCE} = -\log(0.31) = 1.17$

PROPERTIES OF CE LOSS

$$L_{CE} = - \sum_{i=1}^n [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

- Ranges from 0 (perfect predictions) to ∞
Lower the value, better the classifier
- *Cross-entropy* between the true distribution $P(y | \mathbf{x})$ and predicted distribution $P(\hat{y} | \mathbf{x})$

OPTIMIZATION


- We have our **classification function** and **loss function** - how do we find the best w and b ?

$$\theta = [w; b]$$

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{n} \sum_{i=1}^n L_{CE}(y^{(i)}, x^{(i)}; \theta)$$

- Gradient descent:
 - For a differentiable function f :
 - Find direction of steepest slope
 - Move in the opposite direction

GRADIENT DESCENT (1-D)




$$\theta^{t+1} = \theta^t - \eta \frac{d}{d\theta} f(x; \theta)$$

(*f* is differentiable and θ is the w in the figure above)


GRADIENT DESCENT FOR LR

- Cross entropy loss for logistic regression is **convex** (i.e. has only one global minimum)
 - No local minima to get stuck in
- Deep neural networks are not so easy
 - Non-convex




LEARNING RATE

- Updates: $\theta^{t+1} = \theta^t - \eta \frac{d}{d\theta} f(x; \theta)$
- Magnitude of movement along gradient
- Higher/faster learning rate = larger updates to parameters



Too small: converge
very slowly




Too big: overshoot and
even diverge

GRADIENT DESCENT WITH VECTOR WEIGHTS

- In LR: weight w is a vector
- Express slope as a partial derivative of loss w.r.t each weight:

$$\nabla_{\theta} L(f(x; \theta), y) = \begin{bmatrix} \frac{\partial}{\partial w_1} L(f(x; \theta), y) \\ \frac{\partial}{\partial w_2} L(f(x; \theta), y) \\ \vdots \\ \frac{\partial}{\partial w_n} L(f(x; \theta), y) \end{bmatrix}$$



GRADIENT DESCENT WITH VECTOR WEIGHTS

- In LR: weight w is a vector
- Express slope as a partial derivative of loss w.r.t each weight:

$$\nabla_{\theta} L(f(x; \theta), y) = \begin{bmatrix} \frac{\partial}{\partial w_1} L(f(x; \theta), y) \\ \frac{\partial}{\partial w_2} L(f(x; \theta), y) \\ \vdots \\ \frac{\partial}{\partial w_n} L(f(x; \theta), y) \end{bmatrix}$$

- Updates: $\theta(t+1) = \theta(t) - \eta \nabla L(f(x; \theta), y)$

GRADIENT FOR LOGISTIC REGRESSION

- Cross entropy loss:

$$L_{CE} = - \sum_{i=1}^n [y^{(i)} \log \sigma(w \cdot x^{(i)} + b) + (1 - y^{(i)}) \log(1 - \sigma(w \cdot x^{(i)} + b))]$$

- Gradient:

$$\frac{dL_{CE}(w, b)}{dw_j} = \sum_{i=1}^n [\sigma(w \cdot x^{(i)} + b) - y^{(i)}] x_j^{(i)}$$

$$\frac{dL_{CE}(w, b)}{db} = \sum_{i=1}^n [\sigma(w \cdot x^{(i)} + b) - y^{(i)}]$$

- Recall:

$$\frac{d}{dx} \ln(x) = \frac{1}{x} \quad \frac{d\sigma(z)}{dz} = \sigma(z)(1 - \sigma(z))$$

QUIZ: DERIVE THE DERIVATIVE OF CE LOSS

- Given that: $\frac{d}{dx} \ln(x) = \frac{1}{x}$ $\frac{d\sigma(z)}{dz} = \sigma(z)(1 - \sigma(z))$

$$L_{CE} = - \sum_{i=1}^n [y^{(i)} \log \sigma(w \cdot x^{(i)} + b) + (1 - y^{(i)}) \log(1 - \sigma(w \cdot x^{(i)} + b))]$$

Derive (showing steps) that:

$$\frac{dL_{CE}(w, b)}{dw_j} = \sum_{i=1}^n [\sigma(w \cdot x^{(i)} + b) - y^{(i)}] x_j^{(i)}$$

GRADIENT FOR LOGISTIC REGRESSION

- Cross entropy loss:

$$L_{CE} = - \sum_{i=1}^n [y^{(i)} \log \sigma(w \cdot x^{(i)} + b) + (1 - y^{(i)}) \log(1 - \sigma(w \cdot x^{(i)} + b))]$$

- Gradient:

$$\frac{dL_{CE}(w, b)}{dw_j} = \sum_{i=1}^n [\overbrace{\sigma(w \cdot x^{(i)} + b) - y^{(i)}}^{\hat{y}^{(i)}}] x_j^{(i)}$$

input
feature
value

$$\frac{dL_{CE}(w, b)}{db} = \sum_{i=1}^n [\sigma(w \cdot x^{(i)} + b) - y^{(i)}]$$

$$\frac{dL_{CE}(w, b)}{dw_j} = \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)}$$

STOCHASTIC GRADIENT DESCENT


- Online optimization
- Compute loss and minimize after *each training example*

```
function STOCHASTIC GRADIENT DESCENT( $L()$ ,  $f()$ ,  $x$ ,  $y$ ) returns  $\theta$ 
    # where: L is the loss function
    #      f is a function parameterized by  $\theta$ 
    #      x is the set of training inputs  $x^{(1)}$ ,  $x^{(2)}$ , ...,  $x^{(n)}$ 
    #      y is the set of training outputs (labels)  $y^{(1)}$ ,  $y^{(2)}$ , ...,  $y^{(n)}$ 

     $\theta \leftarrow 0$ 
    repeat til done    # see caption
        For each training tuple  $(x^{(i)}, y^{(i)})$  (in random order)
            1. Optional (for reporting):      # How are we doing on this tuple?
                Compute  $\hat{y}^{(i)} = f(x^{(i)}; \theta)$     # What is our estimated output  $\hat{y}$ ?
                Compute the loss  $L(\hat{y}^{(i)}, y^{(i)})$     # How far off is  $\hat{y}^{(i)}$  from the true output  $y^{(i)}$ ?
            2.  $g \leftarrow \nabla_{\theta} L(f(x^{(i)}; \theta), y^{(i)})$     # How should we move  $\theta$  to maximize loss?
            3.  $\theta \leftarrow \theta - \eta g$                       # Go the other way instead
    return  $\theta$ 
```

STOCHASTIC GRADIENT DESCENT

- *Online* optimization
- Compute loss and minimize after *each training example*



REGULARIZATION

- Training objective:

$$\hat{\theta} = \arg \max_{\theta} \sum_{i=1}^n \log P(y^{(i)} | x^{(i)})$$

- This might fit the training set too well! (including noisy features)
- Poor generalization to the unseen test set — ***Overfitting***
- ***Regularization*** helps prevent overfitting:

$$\hat{\theta} = \arg \max_{\theta} \sum_{i=1}^n \log P(y^{(i)} | x^{(i)}) - \alpha R(\theta)$$



penalize large weights

L2 REGULARIZATION

$$R(\theta) = \|\theta\|^2 = \sum_{j=1}^d \theta_j^2$$

- Euclidean distance of weight vector θ from origin
- L2 regularized objective:

$$\hat{\theta} = \arg \max_{\theta} \sum_{i=1}^n \log P(y^{(i)} | x^{(i)}) - \alpha \sum_{j=1}^d \theta_j^2$$

L1 REGULARIZATION


$$R(\theta) = \|\theta\|_1 = \sum_{j=1}^d |\theta_j|$$

- Manhattan distance of weight vector θ from origin
- L1 regularized objective:

$$\hat{\theta} = \arg \max_{\theta} \sum_{i=1}^n \log P(y^{(i)} | x^{(i)}) - \alpha \sum_{j=1}^d |\theta_j|$$

L2 vs L1 REGULARIZATION

- L2 is easier to optimize - simpler derivation
 - L1 is complex since the derivative of $|\theta|$ is not continuous at 0
- L2 leads to many small weights (due to θ^2 term)
 - L1 prefers *sparse* weight vectors with many weights set to 0 (i.e. far fewer features used)



MULTINOMIAL LOGISTIC REGRESSION

- What if we have more than 2 classes? (e.g. Part of speech tagging, Named Entity Recognition, language model!)
- Need to model $P(y = c \mid x), \forall c \in C$
- Generalize **sigmoid** function to **softmax**

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^k e^{z_j}} \quad 1 \leq i \leq k$$

↖ Normalization

SOFTMAX

- Similar to sigmoid, softmax squashes values towards 0 or 1, turning a vector into a probability distribution
- If $z = [0,1,2,3,4]$, then
 - $\text{softmax}(z) = ([0.0117, 0.0317, 0.0861, 0.2341, 0.6364])$
- For multinomial LR,

$$P(y = c \mid x) = \frac{e^{w_c \cdot x + b_c}}{\sum_{j=1}^k e^{w_j \cdot x + b_j}}$$

$$\log P(y = c \mid x) \propto w_c \cdot x + b_c \quad (\text{log-linear!})$$

FEATURES IN MULTINOMIAL LR

- Features need to include both input (x) and class (c)
- Implicit in binary case

Var	Definition	Wt
$f_1(0, x)$	$\begin{cases} 1 & \text{if “!”} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	-4.5
$f_1(+, x)$	$\begin{cases} 1 & \text{if “!”} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	2.6
$f_1(-, x)$	$\begin{cases} 1 & \text{if “!”} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1.3

LEARNING

- Generalize binary CE loss to multinomial CE loss:

$$\begin{aligned} L_{\text{CE}}(\hat{\mathbf{y}}, \mathbf{y}) &= - \sum_{k=1}^K \mathbf{y}_k \log \hat{\mathbf{y}}_k \\ &= -\log \hat{\mathbf{y}}_c, \quad (\text{where } c \text{ is the correct class}) \\ &= -\log \hat{p}(\mathbf{y}_c = 1 | \mathbf{x}) \quad (\text{where } c \text{ is the correct class}) \\ &= -\log \frac{\exp(\mathbf{w}_c \cdot \mathbf{x} + b_c)}{\sum_{j=1}^K \exp(\mathbf{w}_j \cdot \mathbf{x} + b_j)} \quad (c \text{ is the correct class}) \end{aligned}$$

- Gradient:

$$\begin{aligned} \frac{\partial L_{\text{CE}}}{\partial \mathbf{w}_{k,i}} &= -(\mathbf{y}_k - \hat{\mathbf{y}}_k) \mathbf{x}_i \\ &= -(\mathbf{y}_k - p(\mathbf{y}_k = 1 | \mathbf{x})) \mathbf{x}_i \\ &= - \left(\mathbf{y}_k - \frac{\exp(\mathbf{w}_k \cdot \mathbf{x} + b_k)}{\sum_{j=1}^K \exp(\mathbf{w}_j \cdot \mathbf{x} + b_j)} \right) \mathbf{x}_i \end{aligned}$$