



CSE 4392 SPECIAL TOPICS
NATURAL LANGUAGE PROCESSING

Log-linear Models

1

2024 Spring

LAST TIME

- Supervised classification:
 - Document to classify, d
 - Set of classes, $C = \{c_1, c_2, \dots, c_k\}$
- Naive Bayes:

$$\hat{c} = \arg \max_c P(c)P(d|c)$$

LOGISTIC REGRESSION

- Powerful supervised model
- Baseline approach to most NLP tasks
- Connections with neural networks
- Binary (two classes) or multinomial (>2 classes)

DISCRIMINATIVE MODEL

- Logistic Regression is a *discriminative* model
- Naive Bayes is a *generative* model



DISCRIMINATIVE MODEL

- Logistic Regression:

$$\hat{c} = \arg \max_{c \in C} P(c|d)$$

- Naive Bayes:

$$\hat{c} = \arg \max_{c \in C} P(c) P(d|c)$$



QUIZ



- Given that we want to classify an image into either a dog or a cat (no other choices), name the features you would use (can be numerical or categorical).

USING LOGISTIC REGRESSION

○ Inputs:

1. Classification instance in a **feature representation** $[x_1, x_2, \dots, x_d]$
2. **Classification function** to compute \hat{y} using $P(\hat{y} | \mathbf{x})$
3. **Loss function** (for learning)
4. Optimization **algorithm**

○ Train phase:

- Learn the **parameters** of the model to minimize **loss function**

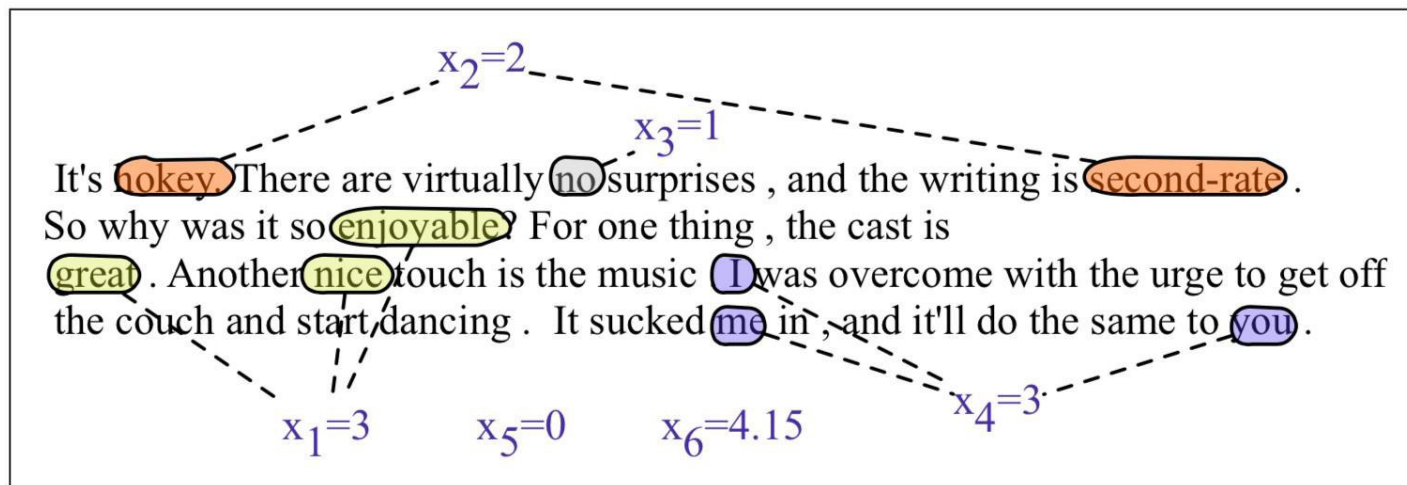
○ Test phase:

- Apply **parameters** to predict class given a new input \mathbf{x}

FEATURE REPRESENTATION

- Input observation: $x^{(i)}$
- Feature vector: $[x_1, x_2, \dots, x_d]$
- Feature j of i^{th} input: $x_j^{(i)}$

SAMPLE FEATURE VECTOR




Var	Definition	Value
x_1	count(positive lexicon) \in doc)	3
x_2	count(negative lexicon) \in doc)	2
x_3	$\begin{cases} 1 & \text{if "no" } \in \text{ doc} \\ 0 & \text{otherwise} \end{cases}$	1
x_4	count(1st and 2nd pronouns \in doc)	3
x_5	$\begin{cases} 1 & \text{if "!" } \in \text{ doc} \\ 0 & \text{otherwise} \end{cases}$	0
x_6	log(word count of doc)	$\ln(64) = 4.15$

CLASSIFICATION FUNCTION

- *Given:* Input feature vector $[x_1, x_2, \dots, x_d]$
- *Output:* $P(y = 1 \mid x)$ and $P(y = 0 \mid x)$ (binary classification)
- Require a *function*, $F : \mathbb{R}^d \rightarrow [0,1]$ (probability)
- Sigmoid (or logistic) Function:

$$\sigma(z) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + \exp(-z)}$$



QUIZ

- Why do we use Sigmoid/Logistic function as our classification function? (Select all that apply)
 - a) Produces a value between 0 and 1
 - b) A partial function with domain $[0, +\infty)$
 - c) Produces a value between -1 to 1
 - d) A total function with domain $(-\infty, \infty)$
 - e) Integrates to 1 from $-\infty$ to ∞
 - f) Differentiable

WEIGHTS AND BIASES

- *Which features are important* and *how much*?
- Learn a vector of **weights** and a **bias**
- **Weights:** Vector of real numbers,

$$\mathbf{w} = [w_1, w_2, \dots, w_d]$$

- **Bias:** Scalar intercept, b
- Given an instance \mathbf{x} :

$$z = \sum_{i=1}^d w_i x_i + b$$

$$\text{or } z = \mathbf{w} \cdot \mathbf{x} + b$$

WHAT IS THE BIAS?

$$z = \mathbf{w} \cdot \mathbf{x} + b$$

- Bias, or intercept, gives the default behavior of the classifier when no useful information about \mathbf{x} is known.
- Try setting x_i to be all 0:

$$z = b$$

- Gives the prior probability distribution of the classes without looking at the input features:

$$\text{prediction_bias} = \text{avg_predictions} - \text{avg of labels in data set}$$

PUTTING IT TOGETHER

- Given \mathbf{x} , compute $z = \mathbf{w} \cdot \mathbf{x} + b$
- Compute probabilities:

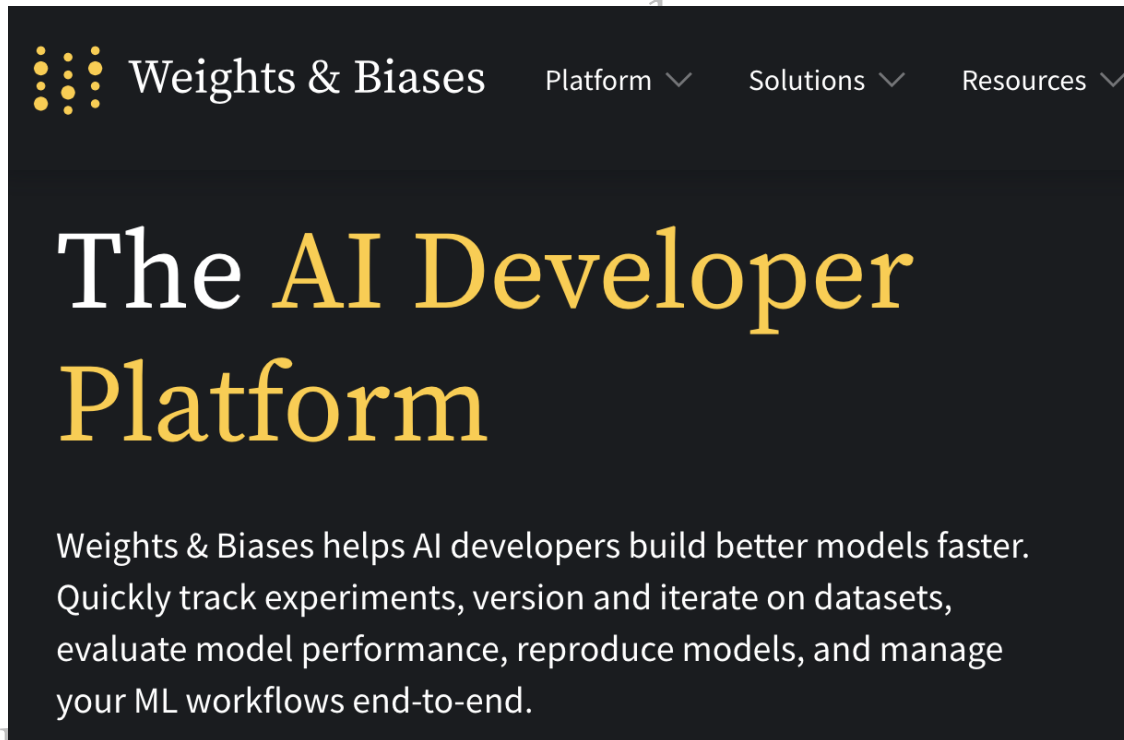
$$\begin{aligned}P(y = 1 | \mathbf{x}) &= \frac{1}{1 + e^{-z}} \\P(y = 1) &= \sigma(\mathbf{w} \cdot \mathbf{x} + b) \\&= \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}} \\P(y = 0) &= 1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b) \\&= 1 - \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}} \\&= \frac{e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}\end{aligned}$$

- Decision boundary:

$$\hat{y} = \begin{cases} 1 & \text{if } P(y = 1 | x) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

PUTTING IT TOGETHER

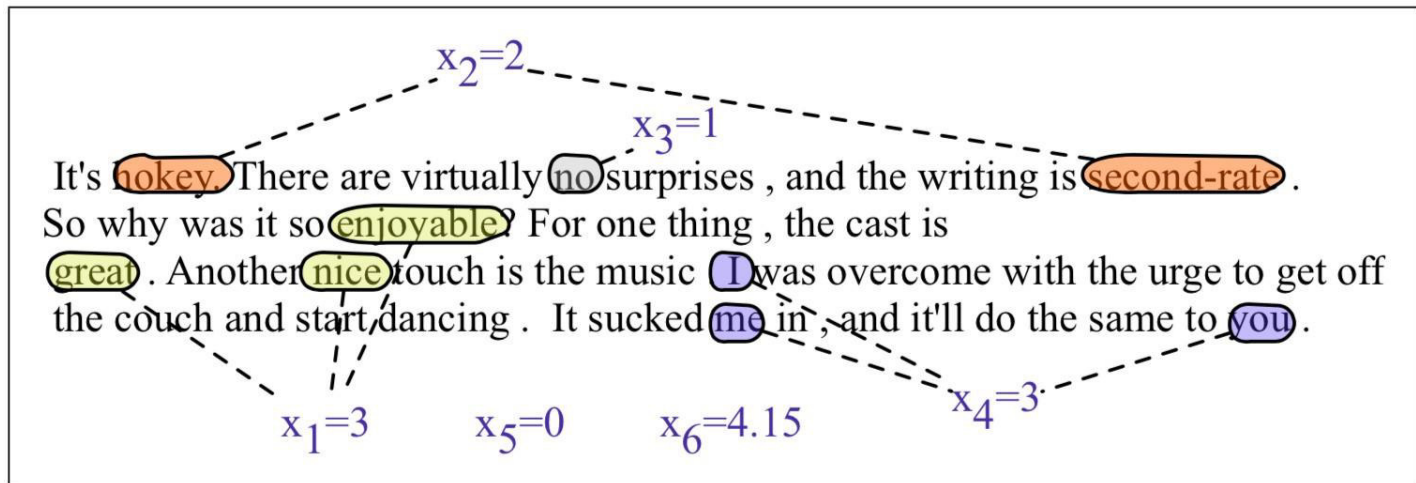
- Given x , compute $z = w \cdot x + b$
- Compute probabilities:



- Decision boundary:

$$\hat{y} = \begin{cases} 1 & \text{if } P(y = 1 | x) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

EXAMPLE: SENTIMENT CLASSIFICATION



Var	Definition	Value
x_1	count(positive lexicon) \in doc	3
x_2	count(negative lexicon) \in doc	2
x_3	$\begin{cases} 1 & \text{if "no" } \in \text{ doc} \\ 0 & \text{otherwise} \end{cases}$	1
x_4	count(1st and 2nd pronouns \in doc)	3
x_5	$\begin{cases} 1 & \text{if "!" } \in \text{ doc} \\ 0 & \text{otherwise} \end{cases}$	0
x_6	log(word count of doc)	$\ln(64) = 4.15$

EXAMPLE: SENTIMENT CLASSIFICATION

Var	Definition	Value
x_1	count(positive lexicon) \in doc)	3
x_2	count(negative lexicon) \in doc)	2
x_3	$\begin{cases} 1 & \text{if "no"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
x_4	count(1st and 2nd pronouns \in doc)	3
x_5	$\begin{cases} 1 & \text{if "!"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
x_6	log(word count of doc)	$\ln(64) = 4.15$

- Assume weights $\mathbf{w} = [2.5, -5.0, -1.2, 0.5, 2.0, 0.7]$ and bias $b = 0.1$

$$\begin{aligned} p(+|x) &= P(Y = 1|x) = \sigma(\mathbf{w} \cdot \mathbf{x} + b) \\ &= \sigma([2.5, -5.0, -1.2, 0.5, 2.0, 0.7] \cdot [3, 2, 1, 3, 0, 4.15] + 0.1) \\ &= \sigma(.805) \\ &= 0.69 \end{aligned}$$

$$\begin{aligned} p(-|x) &= P(Y = 0|x) = 1 - \sigma(\mathbf{w} \cdot \mathbf{x} + b) \\ &= 0.31 \end{aligned}$$

FEATURE DESIGN

- Most important rule: Data is *key*!
- Linguistic intuition (e.g. part of speech tags, parse trees)
- Complex combinations

$$x_1 = \begin{cases} 1 & \text{if “Case}(w_i) = \text{Lower”} \\ 0 & \text{otherwise} \end{cases}$$

$$x_2 = \begin{cases} 1 & \text{if “}w_i \in \text{AcronymDict”} \\ 0 & \text{otherwise} \end{cases}$$

$$x_3 = \begin{cases} 1 & \text{if “}w_i = \text{St. \& Case}(w_{i-1}) = \text{Cap”} \\ 0 & \text{otherwise} \end{cases}$$

- Feature templates
 - Sparse representations, hash only seen features into index
 - Ex. Trigram(*“logistic regression model”*) = Feature #78
- Advanced: Representation learning (we will see this later!)

PROS AND CONS OF LOGISTIC REGRESSION

- More freedom in designing features
 - No strong independence assumptions like Naive Bayes
 - More robust to correlated features (“San Francisco” vs “Boston”) —LR is likely to work better than NB
 - Can even have the same feature twice! (*why?*)
- **However:** Naïve Bayes (NB) often better on very small datasets

LEARNING

- We have our **classification function** - how to assign weights and bias?
- **Goal:** predicted label \hat{y} as close as possible to actual label y
 - **Distance metric/Loss function** between \hat{y} and y :
$$L(\hat{y}, y)$$
 - **Optimization algorithm** for updating weights

LOSS FUNCTION

- Assume $\hat{y} = \sigma(\mathbf{w} \cdot \mathbf{x} + b)$
- $L(\hat{y}, y)$ = Measure of difference between \hat{y} and y . But what form?
- Maximum likelihood estimation (conditional):
 - Choose w and b such that $\log P(y \mid \mathbf{x})$ is maximized for true labels y paired with input x
 - Similar to language models!
 - $\max \log P(w_t \mid w_{t-n}, \dots, w_{t-1})$ given a corpus

CROSS ENTROPY LOSS

- Assume a single data point (x, y) and two classes
- Binary classifier probability (Bernoulli distribution):

$$P(y | x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

- Log probability:

$$\begin{aligned}\log P(y|x) &= \log[\hat{y}^y (1 - \hat{y})^{1-y}] \\ &= y \log \hat{y} + (1 - y) \log(1 - \hat{y})\end{aligned}$$

- CE Loss (we want to minimize):

$$\begin{aligned}-\log P(y|x) &= -y \log \hat{y} - (1 - y) \log(1 - \hat{y}) \\ &= \begin{cases} -\log \hat{y} & \text{if } y = 1 \\ -\log(1 - \hat{y}) & \text{if } y = 0 \end{cases}\end{aligned}$$

CROSS ENTROPY LOSS

- Assume n data points $(x^{(i)}, y^{(i)})$

- Classifier probability:

$$\prod_{i=1}^n P(y | x) = \prod_{i=1}^n \hat{y}^y (1 - \hat{y})^{1-y}$$

(I omitted the (i) here for brevity)

- CE Loss:

$$\begin{aligned} L_{CE} &= -\log \prod_{i=1}^n P(y^{(i)} | x^{(i)}) \\ &= -\sum_{i=1}^n [y^{(i)} \log \hat{y}^{(i)} - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] \end{aligned}$$

EXAMPLE: COMPUTING CE LOSS

Var	Definition	Value
x_1	$\text{count}(\text{positive lexicon}) \in \text{doc}$	3
x_2	$\text{count}(\text{negative lexicon}) \in \text{doc}$	2
x_3	$\begin{cases} 1 & \text{if "no"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
x_4	$\text{count}(\text{1st and 2nd pronouns}) \in \text{doc}$	3
x_5	$\begin{cases} 1 & \text{if "!"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
x_6	$\log(\text{word count of doc})$	$\ln(64) = 4.15$

- Assume weights $w = [2.5, -5.0, -1.2, 0.5, 2.0, 0.7]$ and bias $b = 0.1$
- If $y = 1$ (positive sentiment), $\text{LCE} = -\log(0.69) = 0.37$
- If $y = 0$ (negative sentiment), $\text{LCE} = -\log(0.31) = 1.17$

PROPERTIES OF CE LOSS

$$L_{CE} = - \sum_{i=1}^n [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

- Ranges from 0 (perfect predictions) to ∞
Lower the value, better the classifier
- *Cross-entropy* between the true distribution $P(y \mid \mathbf{x})$ and predicted distribution $P(\hat{y} \mid \mathbf{x})$

OPTIMIZATION


- We have our **classification function** and **loss function** - how do we find the best w and b ?

$$\theta = [w; b]$$

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{n} \sum_{i=1}^n L_{CE}(y^{(i)}, x^{(i)}; \theta)$$

- Gradient descent:
 - For a differentiable function f :
 - Find direction of steepest slope
 - Move in the opposite direction

GRADIENT DESCENT (1-D)



$$\theta^{t+1} = \theta^t - \eta \frac{d}{d\theta} f(x; \theta)$$

(f is differentiable)

GRADIENT DESCENT FOR LR

- Cross entropy loss for logistic regression is **convex** (i.e. has only one global minimum)
 - No local minima to get stuck in
- Deep neural networks are not so easy
 - Non-convex

